

Contextual Conclusion Capture: Conflicting Conclusions Degrade LLM-Judge Discrimination Across Reasoning Domains

Author: Pete Sherratt · Independent researcher · contextual-conclusion-capture@tuta.com

Preprint — 2026-08. All numeric results are produced by preregistered, publicly committed experiments with streamed evidence and an independently implemented re-analysis from the raw rows; see the Reproducibility appendix for exact artifacts, immutable commit references, hashes, and per-result findings files.

Abstract

Large language models are increasingly used to *judge* the outputs of other models, often by comparing a candidate answer against a reference. We identify and characterise a failure mode we call **Contextual Conclusion Capture (CCC)**: an LLM judge's ability to distinguish correct from incorrect candidates deteriorates when a *conflicting conclusion* is present in its evaluation context. Through preregistered experiments across five cost-tier model endpoints, we show that neither an authoritative source label nor an elaborate supporting argument is *necessary* for the effect: a **bare, neutrally-labelled wrong answer is sufficient** to induce capture. We then test whether the failure survives a change of reasoning domain, using three domains that admit an executable oracle for correctness: arithmetic word problems, Python function implementations (graded against a frozen unit-test suite), and relational SQL queries (SQLite). CCC replicates in all three, but its magnitude and *which models are affected* are **domain-dependent**: a judge not significantly captured in the code domain is among the most affected in the SQL domain, where every one of the five models is captured strongly enough to *reverse* its ordering (scoring the wrong answer above the correct one). A separate confirmatory extension to four frontier-tier endpoints reproduces SQL capture for GPT-5.6-sol, Gemini-3.1-Pro and Grok-4.5; finds weaker, model-specific code capture only for GPT-5.6-sol; and finds arithmetic capture only for Claude Fable. Fable's code and SQL contrasts are unmeasurable because provider content filtering is concentrated in injected conditions. A conditional frontier follow-up finds that written verification leaves supported residual capture for Fable arithmetic and GPT code; GPT and Grok SQL have no supported positive residual, while Gemini SQL is unmeasurable because verification responses truncate asymmetrically. A further four-arm model-family extension finds supported SQL capture for Qwen 3.7 Plus, Kimi K2.7 Code, MiniMax M3, and GLM 5.2; arithmetic capture for Qwen, MiniMax, and GLM; and code capture only for MiniMax. We evaluate three mitigations on the cost-tier panel under mirrored correct/wrong-reference sentinels. Prompt-level written verification reduces but does not eliminate capture in any domain. A **hybrid conflict-router** — an

independent model solve, followed by a deterministic comparison of conclusions — recovers discrimination where that comparison is clean. And **context isolation** — never placing the foreign conclusion in the judge's prompt — is the only tested safeguard that **removes the reference-to-judge pathway by construction**, which we verify by showing the judge's prompt is byte-identical whether the reference is correct or wrong. We are careful to distinguish this *structural* guarantee (the pathway is absent) from a behavioural one (the judge is correct): isolation removes the tested channel, not all sources of error. That byte-identical arm then does further work as a **negative control**: its susceptibility is zero by construction, so it measures our estimator rather than the architecture. Read that way it shows our own support rule returning a supported effect on a null contrast in one of four measurable arms — which bounds what small estimates elsewhere in the paper can carry — and it supplies the judge's *unaided* discrimination, against which a wrong reference leaves three of four judges **worse off than no reference at all**. We argue evaluator integrity should be treated as a **per-release, per-domain** property rather than one established once, and release all items, harnesses, preregistrations, and evidence.

1. Introduction

"LLM-as-a-judge" is a default tool for evaluating open-ended model behaviour, powering leaderboards, preference-data pipelines, and regression tests. A judge is typically shown a task, a candidate answer, and sometimes a reference answer, and asked to score or compare. This convenience rests on an unstated assumption: that the judge grounds its verdict in the *evidence* (the task and the candidate), not in whatever other conclusions share its context.

We show that assumption fails in a specific, reproducible way. When a conflicting conclusion — for example, a plausible but wrong "reference" answer — is placed in the judge's context, the judge often stops grading the candidate on its merits and instead grades it by *agreement with the conflicting conclusion*. We call this **Contextual Conclusion Capture (CCC)**.

Contributions:

1. **Mechanism, stated as sufficiency rather than falsification.** A source label ("Solver") and a full supporting argument each add nothing detectable over a neutral bare answer; a bare wrong answer is *sufficient* to induce capture (§5.2). We do not claim authority or persuasion can never matter — only that neither is necessary here.
2. **Generality with heterogeneity.** CCC replicates across three domains with executable correctness oracles — arithmetic, imperative code, relational SQL — and both the size of the effect and *which models are affected* differ by domain (§6). A four-endpoint frontier extension replicates the cross-provider SQL result while narrowing code capture to one judge; four additional model-family arms all reproduce SQL capture. Robustness does not transfer across domains or releases.
3. **What prevents it, and in what sense.** Under mirrored sentinels we compare three mitigations. Written verification is partial in every domain; a hybrid router recovers discrimination where the underlying comparison is clean; and **context isolation removes the reference-to-judge pathway by construction**, which we verify byte-for-byte. We separate this structural claim from any claim of correct judging (§7).

4. **A structural control is also a measuring instrument.** Because the isolated architecture's prompts are byte-identical across reference variants, its susceptibility is zero *by construction* — which makes it a negative control we did not have to design. Reading it as one yields two results the treatment arms cannot: our own support rule returns a supported effect on that null contrast in one of four measurable arms, which bounds what small estimates elsewhere in this paper can carry (§7.1); and the same arm supplies the judge's *unaided* discrimination, against which a wrong reference leaves three of four judges worse off than no reference at all (§7.2). We recommend the design independently of our findings: **run a byte-identical arm alongside the treatment arms, and calibrate the decision rule against it before applying it.**

Every experiment was preregistered before data collection, uses fail-closed handling of missing observations, streams auditable evidence, and was re-analysed by an independently implemented pipeline reading the raw rows. Measuring a judge's integrity is itself a judgement problem, and the study is designed so no model establishes the correctness labels it is evaluated against.

2. Related work

LLM-as-a-judge and its biases. Using strong LLMs as evaluators was popularised by MT-Bench and Chatbot Arena (Zheng et al., 2023) for pairwise preference, and by single-output scoring metrics such as G-Eval (Liu et al., 2023), which prompts the judge with the source document and the candidate and has it generate its own evaluation steps. G-Eval is *reference-free*: the judge is conditioned on auxiliary context, but never on a gold answer. This paper studies the **reference-conditioned** case, in which a gold — or purportedly gold — answer is placed in the judge's prompt alongside the candidate. That is the setting our harness implements throughout, and the conclusion the judge is captured by enters through exactly that slot. Documented biases include position/order effects (Wang et al., 2024) and verbosity / self-preference (Panickssery et al., 2024). CCC is distinct: it is not a preference over surface features of the *candidate*, but a deterioration of *discrimination* caused by a conflicting conclusion elsewhere in context, separable (we show) from authority and argument quality.

Reference conflict, and which way it resolves. The closest work to ours is Lee et al. (2026), who introduce a controlled *swapped-reference* QA framework: the gold answer is replaced with an incorrect entity, candidate answers are aligned to the original and swapped references, and grading reliability is measured under the conflict this induces. They find reliability drops sharply across a broad set of judges, attribute it to the judges' over-reliance on parametric knowledge, and find the failure survives common prompt-based mitigations — which is also what we find, in three different domains and against a different mitigation set (§7).

The two failures resolve in **opposite directions**, and the contrast is the more informative result. In their setting the judge *rejects* the supplied reference: a candidate that correctly matches the swapped reference is graded Incorrect, because the judge's own knowledge of the entity overrides the instruction it was given. In ours the judge *adopts* it: a candidate the oracle certifies as correct is scored below a wrong one, because the conflicting conclusion overrides the evidence in front of it. Domain is the plausible explanation. Entity QA hands the judge a strong

parametric prior to defend; arithmetic, program semantics and relational queries do not, and the conclusion sitting in context fills the vacuum. Consistent with this, our own effect is largest in SQL (§6.2).

Read together the two results are worse than either alone. They do not show that LLM judges systematically over-trust references, nor that they systematically under-trust them. They show that how a judge resolves a reference-belief conflict is **not a stable property of the judge** but a function of the domain it is grading in — which is the case for measuring evaluator integrity per domain and per release (§9) rather than establishing it once.

Sycophancy, anchoring, and belief bias. LLMs tend to agree with assertions in context and defer to stated beliefs (Sharma et al., 2024). CCC can be read as an evaluation-time anchoring effect, but our controls show the anchor need not be authoritative or argued, which distinguishes it from pure authority-deference. The effect also has a much older human analogue: Evans, Barston and Pollard (1983) show that people evaluating syllogisms endorse arguments whose conclusions they believe and reject those they disbelieve, independent of validity, with the bias strongest on *invalid* arguments — the same place a judge's error is most costly. We do not argue the mechanisms are the same. We note that conclusion-driven evaluation is a known failure of reasoners assessed against a formal standard, and that placing a conclusion in an evaluator's context is a design decision, not a neutral one.

Prompt injection. Indirect prompt injection studies untrusted context subverting behaviour (Greshake et al., 2023). CCC is a benign-context analogue: the "injection" is an ordinary wrong reference answer of the kind reference-conditioned evaluation pipelines routinely paste into judge prompts (Lee et al., 2026), and the harm is silent mis-scoring rather than hijacked instructions.

What the family shares, and why CCC needs a different fix. Position bias, verbosity and self-preference, sycophancy, persuasion, authority bias and knowledge-reference conflict are usually catalogued as separate defects. They share one property: **the judge's verdict is a function of something other than the evidence it was asked to grade.** Where they differ is in what that something is, and this is what determines the available defence. Position, verbosity and beauty are properties of the *presentation*, and can be neutralised by randomising or normalising it. Authority and persuasion are properties of the *framing*, and can be attacked by stripping the cue. Self-preference is a property of the *judge*, and can be attacked by changing the judge. Each of those defences works because the offending feature is separable from the content the judge needs.

CCC is the case where it is not. The feature that moves the verdict is the conflicting conclusion itself — not its wrapper, since we show a bare neutral answer suffices (§5.2) — and a conclusion is precisely the payload a reference field exists to carry. There is no cue to strip, no presentation to normalise, and no second judge that would not face the same context. That is why the mitigations in §7 are structural rather than corrective: the only intervention with a guarantee we can verify is the one that changes what enters the prompt, because the defect is not a bias layered on the evidence but a substitution for it.

Reasoning and verification. Chain-of-thought (Wei et al., 2022) and self-consistency (Wang et al., 2023) improve task accuracy. A natural hypothesis is that asking the judge to reason/verify first would cure CCC; we test this and find it reduces but does not eliminate the effect in any domain. The faithfulness literature predicts as much. Turpin et al. (2023) show that a biasing feature in the input can change a model's answer while the chain-of-thought rationalises the new answer without ever mentioning the feature, and Lanham et al. (2023) find faithfulness *decreasing* with scale on most tasks. Chen et al. (2025) extend this to reasoning models, which reveal their use of an inserted hint in under 20% of cases. A written verification step therefore cannot be assumed to expose a captured verdict — it may simply supply a well-formed argument for it, which is what §5.1 observes directly.

Auxiliary information as an attack surface. Closest in spirit to the setting studied here, Li et al. (2025) show that the auxiliary information supplied to improve judges on complex tasks — rubrics, background, and reference answers — is itself exploitable, and report Large Reasoning Models as *paradoxically* more vulnerable. Li et al. (2026) study scoring-based rather than comparative judges, the mode we use throughout, and identify a *reference answer score bias* among biases originating in the scoring prompt rather than in the evaluated output. Koo et al. (2024) benchmark six cognitive biases across sixteen evaluators and find bias in roughly 40% of comparisons. CCC differs from all three in isolating a single factor — the conclusion — and testing it against an executable oracle rather than against human preference.

3. Definitions

Each judging cell yields a score $s \in [0,100]$. For an item with a known-correct candidate and a matched known-wrong candidate, **discrimination** is

$$D = \text{mean } s(\text{correct}) - \text{mean } s(\text{wrong}), \text{ with } s \in [0,100], \text{ so } \mathbf{D} \in [-100, +100].$$

$D > 0$ means the judge scores correct candidates above wrong ones. We introduce a **conflicting conclusion** into the judge's context — a claim, presented as a neutral reference note, that the answer is the *wrong* value — and measure

$$\text{harm} = D(\text{no injection}) - D(\text{injection}), \mathbf{harm} \in [-200, +200].$$

In the mirrored-reference design, the same item is judged with a *correct* and a *wrong* reference, and

$$\text{susceptibility} = D(\text{correct reference}) - D(\text{wrong reference}), \in [-200, +200].$$

The ± 200 range explains estimates above 100: a judge that goes from $D = +90$ (scores correct above wrong) to $D = -90$ (scores wrong above correct) has $\text{harm} = +180$.

Contextual Conclusion Capture is $\text{harm} > 0$: a conflicting conclusion in context degrades a judge's ability to distinguish correct from incorrect candidates. When harm exceeds the baseline D , discrimination becomes negative — the judge **reverses**.

4. Method

Models. All runs were executed through **OpenRouter**; every model identifier below is an OpenRouter slug. The initial panel comprised five hosted endpoints across four providers: `openai/gpt-4o-mini`, `anthropic/claude-haiku-4.5`, `google/gemini-2.5-flash`, `deepseek/deepseek-chat`, `meta-llama/llama-3.3-70b-instruct`. A later confirmatory frontier panel tested `openai/gpt-5.6-sol`, `google/gemini-3.1-pro-preview`, `x-ai/grok-4.5`, and `~anthropic/claude-fable-latest` under the score-only protocol. The final four-arm extension tested `qwen/qwen3.7-plus`, `moonshotai/kimi-k2.7-code`, `minimax/minimax-m3`, and `z-ai/glm-5.2` under endpoint-valid, separately frozen response protocols. Model names are provider aliases whose backing may change over time; run dates and the metadata each run records are in the Reproducibility appendix. Results compare endpoint behaviour, not model quality rankings.

Operational gold under a frozen oracle. Correctness is never decided by a model. In every domain the correct and wrong candidates are produced or graded by an executable oracle: arithmetic computed in code; Python graded against a **frozen unit-test suite** run in a sandbox; SQL results from SQLite. We call the resulting label *operational gold under the frozen oracle* rather than absolute ground truth. This distinction matters most for code: passing a fixed unit-test suite is not a proof of general semantic correctness, and the wrong candidates are hand-authored single-fault variants whose faults were audited to be behaviour-visible on the suite. Arithmetic and SQL have exact executable answers; code correctness is relative to the frozen tests. What all three share, and what the study requires, is that **no model establishes the labels** — avoiding the circularity of measuring a judge with a judge.

Design. Each domain runs a two-stage design. **Stage 1 (injection)** crosses items $\times$ 5 models $\times$ 4 conditions $\times$ 2 candidate types $\times$ 2 protocols $\times$ 3 repetitions. Conditions factor the conflicting conclusion into *content* $\times$ *label*: `no_injection`; `neutral/answer_only` (bare wrong result — the **primary** condition); `neutral/full_wrong_rationale`; and `solver/full_wrong_rationale` (authoritative label). Protocols are `score_only` and `verify_written` (work it out first, then score). **Stage 2 (architectures)** is *conditional*: it runs only on the models admitted by the Stage-1 capture threshold (§5.3, §8), crossing four architectures $\times$ mirrored correct/wrong references $\times$ candidates $\times$ repetitions. Figure 1 shows all three reference conditions on a single judge, and the discrimination measure by which every later result is scored.

Figure 1. Contextual Conclusion Capture and the isolation control




Figure 1. Contextual Conclusion Capture and the isolation control. One judge (GPT-4o mini) over the same 16 confirmatory items. (a) and (b) differ only in whether the exposed reference is correct or wrong; (c) removes it from the prompt entirely. Throughout, $D = \text{mean } s(\text{correct candidate}) - \text{mean } s(\text{wrong candidate})$; “correct” and “wrong” in the panel titles describe the **reference**, which is a separate factor. Under isolation D is 32.5 when the withheld reference was wrong and 32.7 when it was correct, and all 480 isolated judge prompts were byte-identical across reference variants.

Frontier extension. The frontier experiment reran Stage 1's four injection conditions under `score_only` only (Phase 1): 16 arithmetic items, 16 code items, and 24 SQL items × 4 judges × 4 conditions × 2 candidate types × 3 repetitions = 5,376 cells. The confirmatory v3 schedule used seed `305774821`, deterministic `sha256(domain)` ordering, a frozen completeness floor, fail-closed missingness, and an injection-balance safeguard. Conditional Phase 2 then admitted the five measurable domain-judge pairs with supported Phase-1 capture and ran the same design under `verify_written` only: Fable arithmetic, GPT code, and GPT/Gemini/Grok SQL (2,496 cells). Its primary quantity is the **residual** harm under verification; mitigation deltas are secondary. No frontier architecture phase was run.

Four-arm extension. Each of four model arms used the same 16 arithmetic, 16 code, and 24 SQL items × 4 conditions × 2 candidate types × 3 repetitions = 1,344 cells. Compatibility pilots froze provider routing, output contract, budgets, reasoning policy, and transport concurrency before each full run. Qwen used bounded reasoning, Kimi used provider-default native reasoning with sufficient headroom, and MiniMax/GLM used strict one-field JSON with reasoning disabled and zero observed reasoning tokens. These arms are compared descriptively but not pooled as protocol-identical.

Statistics. For each model and contrast, the estimate is a within-item paired difference, aggregated by first averaging a cell's 3 repetitions, then forming the item-level discrimination, then averaging over items. Uncertainty is an **item-clustered nonparametric bootstrap**

(resampling items with replacement): the analysis adapters use $B = 4,000$ resamples; the independent re-analysis uses $B = 4,000\text{--}10,000$; seeds are fixed in code. We report 95% percentile intervals. An effect is **supported** if and only if its interval excludes zero in the predicted direction *and* a completeness floor is met ($\geq 75\%$ of items). **No multiplicity correction is applied**: the primary contrast is a single preregistered test per model; all secondary and mechanism contrasts are reported descriptively with intervals and are **not** equivalence tests — an interval covering zero is not evidence of no effect. Where we compare models, overlapping intervals are treated as a conservative tie, not a significance test.

Preregistration and integrity. For each confirmatory run we froze, before any API call: items and their hash, conditions, primary/secondary contrasts, decision rule, repetitions and schedule seed, model set, and missingness policy. Analysis is **fail-closed**: an item enters a contrast only if all required repetitions of every required cell succeeded; missingness is reported before any estimate and never interpreted as safety. Evidence is streamed to append-only JSONL with a hash-keyed prompt manifest; the deduplication key is the frozen cell identity; at most one successful observation per cell is permitted. In the mechanically-gradable domains the run recomputes a **gold-signature** at start and aborts if the local oracle would produce different results (this fired usefully in SQL when the run machine's SQLite 3.50.4 differed from the 3.45.1 of development, and matched).

5. Establishing the phenomenon and its mechanism (arithmetic)

5.1 The phenomenon, and the fragility of prompt fixes

In a 16-item arithmetic pilot, injecting a poisoned reference collapsed discrimination for every model under score-only judging. A 2×2 factorial separating a *verify instruction* from a *requirement to show written work* found the instruction alone inert, written work the main lever but incomplete for some models, and only the combination driving the measured gap to zero — an interaction, not a main effect. A preregistered wording-sensitivity matrix showed the prompt-level fix is model-specific and phrasing-fragile: one model regressed on the most rigid "structural-looking" template. An audit of failing transcripts found the judge computing the correct answer in its own working and then scoring the correct candidate zero — in one case asserting "8 does not match my calculation of 8." The verdict was controlled by the reference, not by the demonstrably-correct computation. Motivation: **visible reasoning does not reveal what controlled the verdict.**

5.2 Provenance and rationale are not necessary

A provenance $\times$ content factorial, confirmed in a separate preregistered run, isolates the active ingredient. Figure 2 plots it.




Figure 2. The bare conclusion is the active ingredient. Bare-conclusion harm (left) beside the rationale and provenance increments (centre, right) on a shared scale, so the long left-hand bars and the near-zero increments are read against the same ruler. Both increments are nested: the centre panel is rationale-over-bare and the right is solver-over-neutral with rationale held constant. Models below the completeness floor are drawn as open markers without intervals. The figure carries its own equivalence caveat, which applies to every interval in the two right-hand panels.

Naming the source ("Solver") added no detectable capture over a neutral label in any model, and for one model *reduced* it; a full wrong rationale did not add detectable capture over a bare wrong answer. We therefore conclude that **neither authority nor argument is necessary** for CCC in these experiments, and that a bare neutral conclusion is **sufficient**. We do not claim these factors can never matter: the mechanism contrasts are not equivalence tests, and their intervals covering zero bound, but do not exclude, small effects.

This does not contradict the authority-bias literature, and the distinction matters.

Chen et al. (2024) find Authority Bias in both human and LLM judges under a reference-free framework, using a *fake references* perturbation among others, and go on to exploit it as an attack. Hwang et al. (2025) formalise seven persuasion techniques including Authority, embed them in otherwise identical solutions to mathematical problems — a domain where correctness is independent of style — and report inflated scores for incorrect solutions, unmitigated by model size and persisting under counter-prompting. Both results are about what an authority signal *can* do.

Our claim is the complementary one, and it is weaker in scope and stronger in consequence. We do not test whether an authority label can increase capture; we test whether one is *required to produce it*, and find it is not. The practical difference is where each result points. If capture needed an authority cue, the defence would be to strip such cues from evaluation context — a filtering problem, and a tractable one. Because a bare, neutrally-labelled wrong answer is sufficient, there is no cue to filter: the payload is the conclusion itself, which is precisely the content a reference field exists to carry. That is why \$7 pursues structural separation rather than sanitisation. A pipeline that successfully removed every authority marker from its judge prompts would not have addressed the effect measured here.

5.3 Conditional architecture test

A causal architecture experiment (16 items, mirrored references) compared four pipelines: a contaminated score-only judge; a contaminated verify-written judge; a **context-isolated** judge that never receives the reference; and a **hybrid conflict-router** (described in §7). Because Stage 2 runs only on models the Stage-1 threshold admitted, its estimates are **conditional on selection** and its denominators are *admitted measurable models*, not all five endpoints. On this item set, context isolation was the most consistent safeguard and passed a byte-level audit (every isolated prompt hash-identical across reference variants); the router helped some models; written verification attenuated but did not abolish capture. Full numbers:

```
experiments/results/FINDINGS_contextual_conclusion_capture_confirmatory.md,  
.../FINDINGS_contextual_capture_architecture.md.
```

6. Domain generalization

Both prior domains are, at bottom, "compute a deterministic result." To test whether CCC is an artifact of that task shape, we replicated the full two-stage design in two further domains with executable correctness oracles. **We treat domain as an observed factor, not a randomized one:** the domains differ not only in reasoning type but in prompts, candidate construction, baseline task difficulty, answer representation, and score saturation. We therefore report *observed* magnitudes per domain and do not attribute magnitude differences causally to "domain."

6.1 Code (imperative program semantics)

The judge scores a Python implementation against a specification; gold is a frozen unit-test suite run in a sandbox. The bare-conclusion primary was **supported in 4 of 5 models** (+36 to +44 points); DeepSeek's interval included zero by 0.21 points and, per the frozen rule, was called *not supported*. A side result: Claude Haiku, which in arithmetic refused the bare-score format on 42% of score-only cells and was unmeasurable there, complied fully in code (0 failures) and, once measurable, was captured (+40) — its earlier non-compliance was format-specific, not protective. In the conditional Stage 2 (4 admitted models; **DeepSeek was not tested architecturally in code**), context isolation and the router were each supported for 3 of 4 *admitted measurable* models (isolation byte-audited, 384/384 identical prompt pairs); written verification was supported for **no** model. `experiments/results/FINDINGS_ccc_codedomain.md`

6.2 Relational (declarative SQL)

The judge scores a claimed query result against a frozen SQLite fixture; gold is the SQLite oracle, protected by the gold-signature gate. This domain showed the **largest observed capture** of the three. The bare-conclusion primary was supported in **all five models at +106 to +153 points**. A per-arm decomposition shows a *reversal*: with no injection the judge scores correct high and wrong low; with the bare wrong-result note, correct collapses toward 0 and wrong rises toward the injected value, so discrimination goes negative. The same judges discriminate correctly when not injected, ruling out a pure saturation artifact.

Model	baseline D	injected D	correct: base→inj	wrong: base→inj
gpt-4o-mini	+31.6	−80.6	79.9 → 0.0	48.3 → 80.6
claude-haiku-4.5	+54.2	−58.3	79.2 → 4.2	25.0 → 62.5
gemini-2.5-flash	+41.2	−83.3	63.5 → 0.0	22.3 → 83.3
deepseek-chat	+67.8	−85.4	92.1 → 2.1	24.2 → 87.5
llama-3.3-70b	+41.4	−66.7	80.6 → 0.0	37.3 → 66.7

All five entered Stage 2. Isolation drove the correct-vs-wrong-reference gap to ≈ 0 for every model (byte-identical prompts, 720/720). The router recovered discrimination for all five. Written verification's mitigation *delta* is large, but the honest measure — residual capture still present under verification — remains supported for 4 of 5 models (small for GPT and Claude, ≈ 0 only for Gemini, +50.9 for DeepSeek and +56.2 for Llama). [experiments/results/FINDINGS_ccc_sql.md](#)

6.3 Cross-domain synthesis

	Arithmetic	Code	SQL
Primary capture (score-only)	supported (4/4 measurable)	supported 4/5	supported 5/5
Observed magnitude (points)	+39 to +88	+12 to +44	+106 to +153
Isolation (byte-audited)	most consistent	3/4 admitted	5/5, gap → 0
Hybrid router	partial (some models)	3/4 admitted	5/5
Written verification	partial	supported for none	partial (residual 4/5)

Read this table conditionally, not as a like-for-like comparison. Denominators are *admitted, measurable models* — the conditional Stage-2 subset — not all five endpoints; "3/4" excludes both non-admitted and unmeasurable models. Stage 2 ran only on models the Stage-1 capture threshold admitted, so its rows are estimated on a set selected for showing the effect, which can inflate them relative to an unconditional panel. The rows are not exchangeable across columns.

Two findings cut across the table. First, **capture and its severity are model- and domain-dependent**: DeepSeek is *not* significantly captured in code (+10.2, CI includes 0) yet is *the most* captured in SQL (+153). A judge's robustness in one domain does not transfer to another. Second, **safeguard efficacy is also domain-dependent, with one structural invariant**: written verification is partial everywhere; the router recovers where the underlying conclusion-comparison is clean; and **only context isolation removes the reference-to-judge pathway by construction**, which we can audit at the byte level. Its guarantee is about the *pathway*, not about the correctness of the resulting judgement.

We note, without claiming, a pattern consistent with the data: the domain with the largest observed capture (SQL) also has the lowest baseline discrimination, consistent with judges leaning more on an injected result when the task is harder to verify. Domain was not randomized, so this is a hypothesis for future work, not a causal result.

6.4 Confirmatory frontier-tier extension

The preregistered v3 extension asked whether the Stage-1 pattern survives at a stronger model tier. The table reports the primary bare-conclusion harm; **SUPPORTED** requires both a positive 95% interval and the frozen completeness/balance safeguards. Figure 3 plots the same estimates.




Figure 3. Bare-conclusion harm at the frontier tier: four judges × three domains on one shared scale. Each panel carries its own item count and completeness floor, so the fail-closed rule can be checked against the plotted n rather than taken on trust. Four of the twelve cells are supported, three of them SQL. Open markers are cells the frozen balance safeguard reports unmeasurable; the figure states on its face why.

Domain	GPT-5.6-sol	Gemini-3.1-Pro	Grok-4.5	Claude Fable
Arithmetic	not supported (−7.1)	not supported (+0.2)	not supported (−0.5)	SUPPORTED +62.9
Code	SUPPORTED +11.2	not supported (−2.5)	not supported (+1.4)	unmeasurable — content-filtered
SQL	SUPPORTED +50.6	SUPPORTED +40.9	SUPPORTED +27.4	unmeasurable — content-filtered

The strongest cross-tier result is **SQL capture**: GPT, Gemini, and Grok all show substantial harm, with descriptive provenance increments of +88, +74 and +33 points. Code is weaker and model-specific: only GPT is supported. This clean confirmation also rejected a fragile result: Grok's exploratory v2 code estimate was +5.7 with a lower interval bound of +0.31, but v3 shrank it to +1.4 with an interval crossing zero. Arithmetic shows the inverse endpoint pattern, with Fable alone captured strongly and no positive CCC support for the other three judges.

GPT's code result should be read as directional. Its estimate is +11.19 [+4.83, +21.05], and that lower bound sits inside the ± 7.8 -point band which §7.1 derives from a contrast that is null by construction. We have not measured the noise floor of *this* arm — §7.1's band comes from the architecture experiment, on arithmetic items — so we cannot say the two are interchangeable. But the asymmetry cuts one way: we downgraded a claim of our own in §7.1 on exactly this reasoning, and declining to apply it here because the result is one we would rather keep would be the more suspect choice. The SQL cells (+50.6, +40.9, +27.4, lower bounds +23.1, +24.2, +9.7) and Fable's arithmetic (+62.9, lower bound +41.6) are untouched by this; they clear the band by wide margins.

Fable is not interpretable as resistant in code or SQL. Its endpoint returned provider content_filter blocks for 148 code cells and 27 SQL cells, concentrated in injected conditions. That treatment-correlated missingness biases the surviving contrast. Under the preregistered balance safeguard, both domains are therefore **unmeasurable**, including the nominal positive SQL estimate. The content-filter outcome is a property of this endpoint and stimulus combination, not evidence for or against CCC. experiments/results/FINDINGS_ccc_frontier.md

Conditional Phase 2 tested written verification only on the five measurable Phase-1-positive pairs. The preregistered residual-first analysis gives:

Domain / judge	Phase-1 harm	Residual under verify_written [95% CI]	Verdict
Arithmetic · Claude Fable	+62.9	+19.38 [+1.79, +39.13]	residual capture
Code · GPT-5.6-sol	+11.2	+8.10 [+3.02, +15.48]	residual capture
SQL · GPT-5.6-sol	+50.6	−0.13 [−0.38, 0.00]	positive residual not supported
SQL · Gemini-3.1-Pro	+40.9	not estimated	unmeasurable — truncation-skewed
SQL · Grok-4.5	+27.4	+9.51 [−0.21, +22.01]	positive residual not supported

Written verification therefore attenuates capture strongly for Fable arithmetic and for measurable SQL judges, but it is not dependable: Fable retains a supported arithmetic residual, GPT retains a supported code residual, and Grok's SQL interval still permits a +22-point residual. GPT SQL is tightly bounded near zero, but the preregistration does not treat a zero-touching interval as proof of elimination. Gemini SQL fails the 5% balance gate because four primary injected responses truncate, creating a 5.56% completion gap; its nominal survivor estimate is not interpreted. The descriptive paired mitigation delta is large for Fable arithmetic (+42), GPT SQL (+51), and Grok SQL (+18), but not clearly positive for GPT code (+2.8 [−0.2, +5.8]).

experiments/results/FINDINGS_ccc_frontier_phase2.md

6.5 Four-arm model-family extension

Four further full arms used the same primary bare-conclusion contrast. All contain 1,344/1,344 successful unique cells and pass the frozen condition-balance safeguards. Figure 4 shows each arm's clean and injected discrimination as a pair.




Figure 4. Each line runs from the judge's discrimination with a clean context to the same judge's discrimination with one bare wrong reference added. The slope is the harm; what makes it consequential is where the line ends. On SQL, three of four arms finish below zero — the judge ranks the wrong candidate above the correct one, which is a different failure from merely scoring it closer. Note the two arms marked reasoning-off: MiniMax and GLM ran with reasoning disabled and zero observed reasoning tokens, and are captured in SQL at +143.68 and +148.61. Whatever produces capture here does not require the judge to be thinking at length about the task.

Model	Arithmetic	Code	SQL
Qwen 3.7 Plus	SUPPORTED +94.17	not supported (-9.27)	SUPPORTED +165.28
Kimi K2.7 Code	not supported (-11.88)	not supported (+1.46)	SUPPORTED +70.69
MiniMax M3	SUPPORTED +75.17	SUPPORTED +19.65	SUPPORTED +143.68
GLM 5.2	SUPPORTED +52.08	not supported (+13.33)	SUPPORTED +148.61

The common result is SQL: all four arms support capture, despite different providers and response protocols. Arithmetic capture appears in three arms; code capture is supported only for MiniMax. MiniMax and GLM used zero reasoning tokens, so long hidden reasoning is not necessary for the effect. Kimi's separately preregistered native-reasoning arm shows that allowing the endpoint its valid generation protocol does not protect SQL judging. The panel is an endpoint comparison rather than a model-license audit, and protocol-separated effect sizes are not treated as exchangeable. [experiments/results/FINDINGS_ccc_openrouter_openweight.md](#)

7. Mitigations, precisely described

Figure 5 places the three tested safeguards beside the exposed baseline, measured as the reference susceptibility each leaves behind.




Figure 5. Reference susceptibility by architecture. The four panels are ordered as an argument rather than by magnitude: the ordinary setup, then two defences that operate on the same contaminated context, then the one that changes what the judge is given. The subheads carry that distinction, because the rightmost panel differs from the others in kind and not only in size. It is drawn as a control, not a fourth treatment — every true value in it is zero by construction, so its column measures the estimator rather than the architecture, which is the calibration §7.1 rests on.

- **Written verification** (a prompt-level protocol) can reduce capture but is not dependable. In the cost-tier SQL panel, a large mitigation *delta* coexists with supported residual capture for 4 of 5 models. At the frontier tier it leaves supported residuals for Fable arithmetic and GPT code, bounds GPT SQL near zero, leaves a wide zero-crossing interval for Grok SQL, and makes Gemini SQL unmeasurable through treatment-correlated truncation.
- **The conflict-router is hybrid, not fully mechanical.** A model produces an independent conclusion by solving the task in a separate context; deterministic code then canonicalises and compares that conclusion against the external reference; routing (use the exposed path on agreement, quarantine to a fresh reference-free judge on disagreement or unparseable output) is mechanical *after* that comparison. The router therefore inherits the solver's fallibility: in SQL, 334 of 360 solver outputs were parseable, the remainder fail-safing to quarantine. A **deterministic test-oracle router** — using the unit tests or the SQLite oracle itself as the conflict signal, with no model in the loop — is specified but **untested** here, and would be the stronger variant.
- **Context isolation** removes the reference from the judge's prompt entirely. We verify, byte-for-byte, that the judge's prompt is identical whether the reference is correct or wrong (480/480, 384/384, and 720/720 identical pairs in the three architecture experiments), so the reference-to-judge pathway is provably absent. This is a structural property of the pipeline, not a behavioural property of the model: isolation removes the tested channel; it does not guarantee the judge is otherwise correct.

7.1 The isolated arm is a negative control, and it calibrates the support rule

The byte audit has a consequence we should state plainly rather than bank quietly. If the isolated judge's prompt does not change when the reference changes, then its reference susceptibility is *exactly zero by construction* — not small, not approximately zero, but a quantity with no mechanism available to make it nonzero. Every susceptibility we estimate in that arm is therefore a draw from the estimator's own noise, on a contrast whose true value we know. The isolated arm is a **negative control we did not have to design**, and it measures the false-positive rate of the support rule used throughout this paper.

It does not come out clean. Applying the frozen rule — at least 12 of 16 complete items and a 95% item-clustered bootstrap interval excluding zero — to the confirmatory architecture arm:

Model	Isolated susceptibility	Complete items	Rule says
GPT-4o mini	+0.21 [-3.75, +4.38]	16	not supported
Claude Haiku 4.5	+0.00	5	unmeasurable
Gemini 2.5 Flash	+5.42 [+0.62, +11.04]	16	supported — and it cannot be
DeepSeek Chat	-4.38 [-14.17, +7.92]	16	not supported
Llama 3.3 70B	-4.44 [-13.33, +0.00]	15	not supported

One of four measurable arms returns a supported effect on a contrast that is null by construction. We report this as a **false positive of our own procedure**, because that is what it is, and the alternative reading — that a reference absent from the prompt moved the verdict — is excluded by the hash audit.

The cause is identifiable. Each cell is a mean over three repetitions, and the bootstrap resamples *items* while treating those per-cell means as fixed. Within-cell sampling noise from stochastic decoding therefore never enters the interval, and the intervals run narrower than the sampling distribution they are meant to cover. On these rows a null contrast has a per-item standard deviation of **16.0 score points**, so a 16-item mean carries a standard deviation near **4.0 points** from decoding alone. `experiments/check_isolation_null_calibration.py` recomputes this table and that dispersion from the raw rows.

Two things follow, and they point in opposite directions.

The first is a limit on what small estimates in this paper can bear. An effect whose interval clears zero by a few points is inside the region where this arm produces false positives, and should not be read as established on the strength of the rule alone. Applying that filter to the architecture findings, the one supported claim it touches is **Gemini's isolation safeguard gain, +23.44 [+3.65, +44.80]**: its point estimate sits well outside the noise floor but its lower bound does not, so we downgrade it here to *directionally consistent, not independently established*. No other supported effect in this paper has a lower bound inside the band. The capture results the paper is built on — +90.00, +112.50, +129.29 in this same experiment — stand at an order of magnitude above it and are untouched.

The second is a point about the estimand. Isolation's guarantee was never that discrimination stays numerically fixed; a judge sampled at nonzero temperature will move a little between any two runs. The guarantee is that **nothing about the reference can be what moves it**, and that is a claim about the prompt bytes, which we can and do check exhaustively. The +5.42 is the size of the wobble you get for free from decoding. The contaminated arm's +39.90 for the same model on the same items is what a reference pathway looks like when it exists. That the two are separable by nearly an order of magnitude is the isolation result, stated more carefully than "isolation works".

A study that ships a structural null and then reads it honestly is in a better position than one that never had a null to read. We recommend the design be adopted for its own sake: **run a byte-identical arm alongside the treatment arms, and calibrate the decision rule against it before applying it.**

7.2 A wrong reference is worse than no reference

The isolated arm serves a second purpose: because the reference never enters the prompt, it is also a **no-reference baseline** — the judge's discrimination on the same items, unaided.

Reporting the three conditions as levels rather than as differences:

Model	No reference	+ correct reference	+ wrong reference
GPT-4o mini	+32.5	+63.4	−26.6
Gemini 2.5 Flash	+42.7	+59.2	+19.3
DeepSeek Chat	+37.7	+69.7	−14.6
Llama 3.3 70B	+41.7	+68.8	−43.8

All $n = 16$; Claude is below the completeness floor in both arms and is omitted. Paired within item against the no-reference baseline, a **correct** reference helps every judge — +30.94 [+9.58, +53.02] for GPT, +31.98 [+10.63, +57.50] for DeepSeek, +27.08 [+8.33, +50.00] for Llama, and +16.46 [+3.54, +32.60] for Gemini. Reference-conditioned judging earns its place; that is not in dispute.

A **wrong** reference does not simply forfeit that gain. For three of the four measurable judges the verdict lands below where it would have been with no reference at all: −59.06 [−89.27, −29.06] for GPT, −52.29 [−66.88, −37.29] for DeepSeek, and −85.42 [−110.42, −62.50] for Llama. Gemini's estimate runs the same way, −23.44 [−44.79, −4.17], but its lower bound falls inside the noise band of \$7.1, so we read it as directional only — as we do its correct-reference gain above, whose lower bound of +3.54 is inside the same band.

The practical reading is narrow and does not require a mechanism. Supplying a reference to a judge is not a neutral act with an upside and no downside. On these items and these judges, the loss from a wrong reference is larger than the gain from a correct one, and large enough to put the judge below its own unaided performance. A pipeline that cannot bound the error rate of its reference answers is not choosing between a good reference and a slightly worse one; it is

choosing between a good reference and an outcome worse than passing no reference at all.
`experiments/check_no_reference_baseline.py` recomputes every estimate above from the raw rows.

We do not claim to know why the judges defer. The isolated column shows they are not incapable of the task, but what governs the trade — the relative difficulty of recomputing the answer, the judge's confidence in its own solution, or something else — is not identified by this design.

8. Missingness

Missingness is factor-correlated and scientifically informative, so we report it in the main text. Failures include unparseable, truncated, non-compliant, provider-filtered, and worker-level responses; all are retained as evidence, none imputed, and affected items are dropped fail-closed (never counted as safe).

Domain · stage	Cells	Failures	Dominant factor cells
Arithmetic · injection (confirmatory)	3,840	218	claude-haiku × score_only 162/384 (42%); gemini × verify_written 56/384
Arithmetic · architecture	3,840	207	claude-haiku ≈159 (score-only format); gemini ≈47; llama 1
Code · injection	3,840	46	gemini × verify_written 45/384; llama × score_only 1
Code · architecture	3,072	29	gemini × contaminated_verify_written 22/192; gemini × conflict_router 7/192
SQL · injection	5,760	23	gemini × verify_written 18/576; llama × verify_written 4; llama × score_only 1
SQL · architecture	5,760	13	gemini × contaminated_verify_written 5/288; llama across four architectures 8
Frontier · arithmetic	1,536	8	gpt 7 (<code>worker:KeyError: 'choices'</code>); gemini 1 truncated
Frontier · code	1,536	154	Fable 148 content-filtered; gpt 6 <code>worker:KeyError</code>
Frontier · SQL	2,304	30	Fable 27 content-filtered; gemini 2; gpt 1 <code>worker:KeyError</code>
Frontier Phase 2 · arithmetic	384	5	Fable worker HTTP 402: baseline 3, solver 2; primary balance passes
Frontier Phase 2 · code	384	0	GPT complete
Frontier Phase 2 · SQL	1,728	10	Gemini truncation only; primary completion gap 5.56% → unmeasurable
OpenRouter · Qwen full arm	1,344	0	complete under bounded-reasoning protocol
OpenRouter · Kimi full arm	1,344	0	two first attempts recovered; no final missing cells
OpenRouter · MiniMax full arm	1,344	0	complete; zero observed reasoning tokens
OpenRouter · GLM full arm	1,344	0	complete; zero observed reasoning tokens

Two patterns recur and shape interpretation. (i) **Claude Haiku's arithmetic score-only non-compliance** (42%) made it unmeasurable in that domain-stage; it complied fully in code and SQL, so its arithmetic gap reflects output-format behaviour, not safety. (ii) **Gemini (and to a lesser extent Llama) fail on `verify_written`** across domains — long derivations truncated before the score JSON — which is why several verification and mitigation contrasts run below full item counts and one (SQL Gemini mitigation) at $n = 22$. Separately, the SQL router's solver produced **334/360** parseable conclusions; the 26 unparseable ones fail-safe to quarantine,

which is the designed behaviour but also a source of router fallibility. In the frontier run, Fable's filtering is injection-skewed and makes its code/SQL contrasts unmeasurable. Fourteen frontier cells failed before a usable provider choice was recorded (`worker:KeyError: 'choices'`); they were retained and dropped fail-closed. Their low, non-systematic rate changed no support call, although GPT arithmetic finished exactly at the item floor. In frontier Phase 2, five late Fable arithmetic failures were billing errors rather than endpoint content filters and remained below the primary balance limit. Gemini's ten SQL failures were instead long `verify_written` derivations truncated after both attempts; four fall in the primary injected cell, pushing its completion gap just above the frozen limit.

Treatment-correlated refusal is a silent failure mode of API-mediated evaluation, and it deserves naming separately from the statistics. Fable's content filter did not fire at random: it fired overwhelmingly in the injected conditions — the very cells in which capture, if present, would show. A pipeline that simply averaged the returned verdicts would read that as evidence of robustness, because the compromised cells are the ones missing. The filter is not a neutral dropout mechanism; conditioning on "the provider returned a verdict" is conditioning on a variable that the treatment influences, which breaks the contrast rather than thinning it. This generalises past our study: **any evaluation harness that treats a provider refusal as a skipped row, rather than as a fact about the row, will report safety it has not measured.** Our balance safeguard exists for this reason, and it is why the honest output for those cells is *unmeasurable* — a refusal to conclude, not a null.

9. Implications for evaluation design

- **Do not paste a reference answer into the judge's context and rely on a prompt to keep it honest.** Reference-anchoring is real, prompt-level verification is incomplete in every domain, and its efficacy is model- and wording-specific.
- **Do not certify a judge once.** Because capture and safeguard efficacy are domain-dependent, an evaluator validated on one task family carries no guarantee on another. Integrity should be a **per-release, per-domain** property, tested with mirrored correct/wrong-reference sentinels run in the same conditions as the real items.
- **Prefer structural separation to behavioural instruction.** Keep the foreign conclusion out of the first-pass judge context (isolation is the only mitigation whose guarantee we can audit at the byte level); where a mechanical comparison of conclusions is available, add a router — ideally the deterministic-oracle variant — as an escalation layer.

10. Limitations

- **Small, hand-authored micro-domains** (16–24 frozen items each). The study establishes the existence, mechanism, and structure of CCC and the relative ordering of mitigations; it is not a benchmark of any deployed system, and confidence intervals over ≤ 24 items are wide.
- **Conditional Stage-2 selection.** Architecture results are estimated only on models the Stage-1 threshold admitted; DeepSeek was not tested architecturally in code; "k/4" or "k/5" denominators refer to admitted measurable models. Selection on a Stage-1 effect can inflate conditional Stage-2 estimates for the admitted subset; we report them as conditional.

- **Domain is observed, not randomized.** Magnitude differences across domains are confounded with prompt wording, candidate construction, difficulty, answer representation, and saturation; "SQL showed the largest observed capture" is supported, "SQL causes more capture" is not.
- **Operational, not absolute, gold.** Especially for code, correctness is defined relative to a frozen unit-test suite and hand-audited faults, not general semantic correctness.
- **Specific, time-stamped endpoints rather than model classes.** The frontier extension covers four endpoint aliases on 2026-07-20, two of them moving `-latest`/preview routes. It supports no claim about frontier models in general. Fable's treatment-correlated content filtering prevents code/SQL estimation and illustrates why missing output cannot be interpreted as robustness.
- **Protocol-separated four-arm panel.** Kimi requires native-reasoning headroom, Qwen uses bounded reasoning, and MiniMax/GLM use a terse no-reasoning contract. Their results establish endpoint-level replication but are not a license audit and must not be naively pooled as one homogeneous panel.
- **Only one frontier mitigation was tested.** Conditional Phase 2 tests written verification on the five measurable Phase-1-positive pairs. It does not test context isolation or the hybrid router at the frontier tier, and conditional selection limits inference to those admitted pairs.
- **Phase-2 metadata exception.** The three domain blocks reused one output prefix, so the shared meta file was overwritten and ultimately describes code only. The complete domain-specific rows and prompt manifests support independent recomputation, but run configuration for arithmetic/SQL is reconstructed from the preregistration, terminal record, and raw evidence rather than preserved in separate final metadata files.
- **Coarse scores** (frequent 0/100 saturation); key results (the reversal, isolation neutrality) rest on decomposition and preregistered support calls, not on treating point magnitudes as precise.
- **The support rule is anticonservative for small effects.** The item-clustered bootstrap resamples items but holds each cell's three-repetition mean fixed, so decoding noise is outside the interval. Measured against the byte-identical isolated arm, where the true effect is zero by construction, the rule returns a supported effect in one of four measurable cases (§7.1). Estimates whose lower bound falls within a few points of zero should be read as directional. The paper's headline effects are an order of magnitude clear of that band; the one claim we downgrade on these grounds is named in §7.1.
- **Mechanical-gold domains only.** Open-ended judging — where the router's comparison would itself become a judgement — is out of scope, because rigour there would reintroduce the evaluator-for-the-evaluator circularity this method avoids. The deterministic test-oracle router is specified but untested.

11. Conclusion

A conflicting conclusion is *sufficient* to degrade LLM-judge discrimination — it needs neither authority nor argument — and it does so across three domains with executable correctness oracles. Susceptibility is model- and domain-dependent: robustness earned in one setting does not carry to another. The frontier extension reinforces rather than erases that heterogeneity:

cross-provider SQL capture replicates, code capture narrows to one judge, and one endpoint is measurable and captured only in arithmetic. The four-arm model-family extension strengthens the SQL result again: every complete arm is captured in SQL, while arithmetic and code remain heterogeneous. The conditional frontier follow-up also shows that prompt-level verification can sharply reduce harm without becoming a reliable safeguard: residual capture survives for Fable arithmetic and GPT code, and another endpoint becomes unmeasurable through truncation. Among the mitigations tested, prompt-level verification is never sufficient; a hybrid router helps where conclusions can be compared cleanly; and the one intervention whose guarantee we can verify byte-for-byte is the simplest — do not let the foreign conclusion into the judge's context. That removes the tested pathway by construction; it does not by itself make the judge correct. Evaluator integrity is not a property to assume; it is one to gate on, every release and every domain.

Author contributions and AI use

All preregistrations, decision rules, completeness floors, and interpretations were fixed by the author, who takes responsibility for the work as published. Implementation and analysis used AI systems throughout, described here at the level of detail a reader would need to judge what the tooling could and could not have influenced.

Claude (Anthropic) implemented the experimental harnesses, the independent re-analysis pipeline used to recompute every headline contrast from the raw observation rows, and the figure set with its validation checks.

ChatGPT / Codex (OpenAI) produced the work on the `codex/*` branches: the model-family extension runs, the reconciliation of the publication base, and the evidence-audit tooling behind the integrity summary.

Gemini (Google) served as adversarial reviewer. Its frame-theft attack against an earlier design is what forced the move to process isolation, and so is upstream of the architecture result in §7.

None of these systems is listed as an author. Authorship is an accountability claim — that a named party approved the final version and can answer for it — and a language model cannot make that claim. The contribution was real; the attribution belongs here rather than on the byline.

References

Every entry below has been checked against its official proceedings record; venue, year, pages, DOI and arXiv identifier are taken from that record rather than reconstructed. In-text citations use venue years throughout.

1. L. Zheng, W.-L. Chiang, Y. Sheng, S. Zhuang, Z. Wu, Y. Zhuang, Z. Lin, Z. Li, D. Li, E. P. Xing, H. Zhang, J. E. Gonzalez, I. Stoica. **Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena**. Advances in Neural Information Processing Systems 36 (NeurIPS 2023), Datasets and Benchmarks Track (arXiv:2306.05685).

2. Y. Liu, D. Iter, Y. Xu, S. Wang, R. Xu, C. Zhu. **G-Eval: NLG Evaluation using GPT-4 with Better Human Alignment.** Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP), pages 2511–2522. DOI 10.18653/v1/2023.emnlp-main.153 (arXiv:2303.16634).
3. P. Wang, L. Li, L. Chen, Z. Cai, D. Zhu, B. Lin, Y. Cao, L. Kong, Q. Liu, T. Liu, Z. Sui. **Large Language Models are not Fair Evaluators.** Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL 2024), Volume 1: Long Papers, pages 9440–9450. DOI 10.18653/v1/2024.acl-long.511 (arXiv:2305.17926).
4. A. Panickssery, S. R. Bowman, S. Feng. **LLM Evaluators Recognize and Favor Their Own Generations.** Advances in Neural Information Processing Systems 37 (NeurIPS 2024). DOI 10.52202/079017-2197 (arXiv:2404.13076).
5. M. Sharma, M. Tong, T. Korbak, D. Duvenaud, A. Askeel, S. R. Bowman, N. Cheng, E. Durmus, Z. Hatfield-Dodds, S. R. Johnston, S. Kravec, T. Maxwell, S. McCandlish, K. Ndousse, O. Rausch, N. Schiefer, D. Yan, M. Zhang, E. Perez. **Towards Understanding Sycophancy in Language Models.** International Conference on Learning Representations (ICLR 2024) (arXiv:2310.13548).
6. K. Greshake, S. Abdelnabi, S. Mishra, C. Endres, T. Holz, M. Fritz. **Not What You've Signed Up For: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection.** Proceedings of the 16th ACM Workshop on Artificial Intelligence and Security (AISeC '23), pages 79–90. DOI 10.1145/3605764.3623985 (arXiv:2302.12173).
7. J. Wei, X. Wang, D. Schuurmans, M. Bosma, B. Ichter, F. Xia, E. H. Chi, Q. V. Le, D. Zhou. **Chain-of-Thought Prompting Elicits Reasoning in Large Language Models.** Advances in Neural Information Processing Systems 35 (NeurIPS 2022) (arXiv:2201.11903).
8. X. Wang, J. Wei, D. Schuurmans, Q. V. Le, E. H. Chi, S. Narang, A. Chowdhery, D. Zhou. **Self-Consistency Improves Chain-of-Thought Reasoning in Language Models.** International Conference on Learning Representations (ICLR 2023) (arXiv:2203.11171).
9. D. Lee, Y. Hwang, T. Kang, M. Lee, Y. Chae, K. Jung. **Judging Against the Reference: Uncovering Knowledge-Driven Failures in LLM-Judges on QA Evaluation.** Seoul National University and LG AI Research, 2026. arXiv:2601.07506.
10. G. H. Chen, S. Chen, Z. Liu, F. Jiang, B. Wang. **Humans or LLMs as the Judge? A Study on Judgement Bias.** Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP 2024), pages 8301–8327.
11. Y. Hwang, D. Lee, T. Kang, Y. Kim, K. Jung. **Can You Trick the Grader? Adversarial Persuasion of LLM Judges.** Findings of the Association for Computational Linguistics: EMNLP 2025, pages 14632–14651.
12. M. Turpin, J. Michael, E. Perez, S. R. Bowman. **Language Models Don't Always Say What They Think: Unfaithful Explanations in Chain-of-Thought Prompting.** Advances in Neural Information Processing Systems 36 (NeurIPS 2023).
13. T. Lanham, A. Chen, A. Radhakrishnan, B. Steiner, C. Denison, D. Hernandez, D. Li, E. Durmus, E. Hubinger, J. Kernion, et al. **Measuring Faithfulness in Chain-of-Thought Reasoning.** Anthropic, 2023. arXiv:2307.13702.

4. Y. Chen, J. Benton, A. Radhakrishnan, J. Uesato, C. Denison, J. Schulman, A. Somani, P. Hase, M. Wagner, F. Roger, V. Mikulik, S. R. Bowman, J. Leike, J. Kaplan, E. Perez. **Reasoning Models Don't Always Say What They Think**. Alignment Science Team, Anthropic, 2025. arXiv:2505.05410.
5. W. Li, X. Wang, S. Yuan, R. Xu, J. Chen, Q. Dong, Y. Xiao, D. Yang. **Curse of Knowledge: When Complex Evaluation Context Benefits yet Biases LLM Judges**. Findings of the Association for Computational Linguistics: EMNLP 2025, pages 14900–14924.
6. Q. Li, S. Dou, K. Shao, C. Chen, H. Hu. **Evaluating Scoring Bias in LLM-as-a-Judge**. arXiv:2506.22316v4 [cs.CL], 3 February 2026.
7. R. Koo, M. Lee, V. Raheja, J. Park, Z. M. Kim, D. Kang. **Benchmarking Cognitive Biases in Large Language Models as Evaluators**. Findings of the Association for Computational Linguistics: ACL 2024, pages 517–545.
8. J. St. B. T. Evans, J. L. Barston, P. Pollard. **On the conflict between logic and belief in syllogistic reasoning**. Memory & Cognition, 1983, 11(3), pages 295–306.

Reproducibility appendix

"**Independent re-analysis**" here means an analysis pipeline implemented separately from the run adapters, reading the raw streamed rows and recomputing every headline contrast; it does **not** imply a separate human investigator. Both the adapter output and the independent recomputation are in the repository, and they matched.

What can be reproduced

We distinguish two targets that are sometimes both called “reproduction”:

1. **Exact computational re-analysis:** recompute the reported estimates, intervals, missingness decisions, and integrity checks from the committed raw JSONL. This is deterministic, makes no API calls, needs no credential, and should match the tables in §§5–8 to the displayed precision.
2. **Fresh endpoint replication:** send the frozen prompts to newly resolved endpoints and obtain a new sample. This requires an API credential and incurs cost. Because hosted aliases, provider implementations, content filters, and routing can change, a fresh run is a replication of the protocol and estimand—not a promise of byte-identical responses or identical effect sizes.

The offline path is the minimum required check for every reported result. A successful reproduction must (i) exit without an exception, (ii) report no duplicate successful cells, (iii) apply the stated completeness and balance gates before estimates, and (iv) reproduce the corresponding paper values to rounding. `AUDIT: PASS` is additionally required where an executable structural audit is supplied.

Exact offline re-analysis (no API calls)

Clone the repository and work from immutable commits rather than a moving branch:

```
git clone https://github.com/petesherratt-collab/The-Generated-Trace-Leak-Harness.git
cd The-Generated-Trace-Leak-Harness
git checkout 3bb12a8
```

The original arithmetic, code, SQL, and conditional architecture results are reproduced with:

```
python experiments/run_provenance_injection.py --confirmatory --analyse-only
python experiments/analyze_confirmatory_choice_probability.py
python experiments/run_architecture_capture.py --analyse-only
python experiments/analyze_architecture_capture.py
python experiments/run_ccc_codedomain.py --analyse-only
python experiments/run_ccc_codedomain_stage2.py --analyse-only
python experiments/run_ccc_sql.py --analyse-only
python experiments/run_ccc_sql_stage2.py --analyse-only
```

The reconciled publication runner reads the historical Frontier namespaces without weakening current-run isolation. It requires each evidence file to contain either uniformly legacy rows with no `run_id`, or uniformly namespaced rows with one matching non-null `run_id`; mixed or multiple run IDs fail closed. For legacy Phase 2 it infers the admitted models and item set independently per domain, avoiding the disclosed shared-metadata overwrite:

```
python experiments/run_ccc_frontier.py --analyse-only
python experiments/run_ccc_frontier.py --analyse-only --domains arith,code,sql --protocols verify_written
```

The evidence-era runners at `8611bec` (v3) and `9f7437a` (Phase 2) remain independent historical reproduction points. Commit `3bb12a8` is the reviewed unified implementation and adds regression tests for uniform legacy, uniform current, mixed, mismatched, and per-domain legacy Phase-2 cases.

Reanalyse each protocol-separated OpenRouter arm:

```
python experiments/run_ccc_frontier.py --analyse-only --domains arith,code,sql --protocols score_only --ev
python experiments/run_ccc_frontier.py --analyse-only --domains arith,code,sql --protocols score_only --ev
python experiments/run_ccc_frontier.py --analyse-only --domains arith,code,sql --protocols score_only --ev
python experiments/run_ccc_frontier.py --analyse-only --domains arith,code,sql --protocols score_only --ev
```

At the publication snapshot, the independent raw-evidence audit accepts both the legacy Qwen metadata and the later full-arm schema. Run it once per arm with the frozen reasoning ceiling:

```
python experiments/audit_ccc_openweight_evidence.py experiments/results/ccc_openrouter_v1 --prefix ccc_ope
python experiments/audit_ccc_openweight_evidence.py experiments/results/ccc_openrouter_kimi_native_v1 --pr
python experiments/audit_ccc_openweight_evidence.py experiments/results/ccc_openrouter_minimax_m3_v2 --pre
python experiments/audit_ccc_openweight_evidence.py experiments/results/ccc_openrouter_glm52_v2 --prefix c
python -m unittest experiments.test_run_ccc_frontier
```

The last command must report **24 passing tests**. The four audits must end in `AUDIT: PASS`. The analysis commands print the estimates from raw observations; the retained `_audit.txt` and `FINDINGS_.md` files provide line-by-line comparison targets. Offline analysis uses only the Python standard library. The commands above were verified under CPython 3.14.4; the live release gates are stricter and require CPython 3.10–3.13 because those are the runtimes on which the executable gold was cross-checked.

The revision analyses, figures, and the document itself

The material in §7.1, §7.2 and the figures was added after the runs and reads the same committed rows. Nothing in it required a new API call, so all of it belongs to the offline path above.

The two control readings. Both recompute their tables from `architecture_obs.jsonl` and print them; neither transcribes a value from the text:

```
python experiments/check_isolation_null_calibration.py      # §7.1 – the null contrast and the false-posit
python experiments/check_no_reference_baseline.py           # §7.2 – wrong reference vs no reference at al
```

The figures. All seven are regenerated from the raw `*_obs.jsonl` rows by one driver, which re-derives every plotted estimate and checks it against the published value. It **exits non-zero on any mismatch**, so a figure cannot silently drift away from the findings it illustrates:

```
python experiments/make_all_ccc_figures.py
```

Each script also prints its own validation block. The architecture figure additionally re-derives the byte audit from the stored `prompt_sha256` values rather than quoting it: 480 / 480 isolated prompt pairs hash-identical across reference variants, 0 differing. `experiments/FIGURES.md` records what each figure claims, which evidence file it reads, and the drawing conventions the set holds to.

Bootstrap seeds. Every interval in this paper — figures and control readings alike — uses `B = 4,000` and the single seed `20260714`. An earlier revision ran the check scripts on their own seeds, which produced a second, equally valid interval for a quantity already published with another; the estimates agreed and only the resample differed. One seed now covers all of them, so any two reported intervals for the same contrast are identical rather than merely compatible.

The document. The PDF is built from the markdown source by a mechanical converter — no word of the source is altered — which inlines the figures and prints through headless Chromium:

```
paper/build_pdf.sh                                # -> paper/Contextual_Conclusion_Capture.pdf
```

References. `paper/REFERENCES_verified.md` records how each of the eighteen entries was confirmed against its official proceedings record, which of the author's source PDFs supports which claim, and — because the point of that file is the distinction — a correction noting that four entries in its first version were written from recall while presented as checked.

Fresh endpoint replication (API calls and cost)

Fresh calls should be made from the last pre-evidence instrument commit, so the historical evidence cannot be overwritten. Set `OPENROUTER_API_KEY` (or `OPENROUTER_ENV_FILE`) without printing it, run the dry-run/wiring checks first, inspect current pricing and alias resolution, and only then authorize the live command. The relevant clean starting commits and entry points are:

block	pre-evidence commit	live entry point
arithmetic confirmatory	<code>77ef756</code>	<code>python experiments/run_provenance_injection.py --confirmatory</code>
arithmetic architectures	<code>ae422bd</code>	<code>python experiments/run_architecture_capture.py</code>
code Stage 1 / Stage 2	<code>58d4465</code> / <code>824f042</code>	<code>python experiments/run_ccc_codedomain.py --run</code> / <code>python experiments/run_ccc_codedomain_stage2.py --run</code>
SQL Stage 1 / Stage 2	<code>c045b43</code> / <code>fa6eb25</code>	<code>python experiments/run_ccc_sql.py --run</code> / <code>python experiments/run_ccc_sql_stage2.py --run</code>
frontier v3	<code>982b97e</code>	commands frozen in <code>experiments/PREREG_ccc_frontier_v3.md</code>
frontier Phase 2	<code>a57e85a</code>	the three admitted-set commands in <code>experiments/PREREG_ccc_frontier_phase2.md</code>
OpenRouter four-arm extension	<code>7e00b4d</code>	the PowerShell launchers below

For the four OpenRouter arms on Windows, use a supported interpreter and a reviewed credential file:

```
$Python = 'C:\path\to\python3.13.exe'
$EnvFile = 'C:\path\to\.env'
```

```
powershell.exe -NoProfile -ExecutionPolicy Bypass -File .\experiments\run_ccc_openweight.ps1 -Mode Run -Ju
powershell.exe -NoProfile -ExecutionPolicy Bypass -File .\experiments\run_ccc_kimi_native.ps1 -Mode Run -F
powershell.exe -NoProfile -ExecutionPolicy Bypass -File .\experiments\run_ccc_openweight_full_v2.ps1 -Mode
powershell.exe -NoProfile -ExecutionPolicy Bypass -File .\experiments\run_ccc_openweight_full_v2.ps1 -Mode
```

Replace `Run` with `DryRun`, `WiringCheck`, and then `CheckModels` (the last with `-ApproveApiCalls`) before spending on a full arm. Never run these launchers from a checkout already containing the corresponding evidence namespace: they are designed to refuse overwriting finalized evidence. A new replication must retain its own date, resolved model/provider identities, completion metadata, prompts, raw observations, and audit output. It should be reported as a new endpoint-time replication even when every configuration field matches the original.

Immutable references. Cite the evidence-bearing commits, not the mutable branch: code Stage 1 at commit `1309d78`, code Stage 2 at `c850083`; SQL Stage 1 at `26354f8`, SQL Stage 2 at `4581589`. (Frontier v3 evidence is at `8611bec`; its independent audit and findings are at `97790ff`. Frontier Phase 2 evidence is at `9f7437a`; its audit and findings are at `3ef311b`.) (The OpenRouter runner and audit tooling are at `7e00b4d`; preregistrations and compatibility pilots are at `841e9d5`; the four full arms, audits, and combined findings are at `80c0dbe`.) (The fail-closed legacy/current reconciliation, complete Frontier Phase-2 re-analysis, legacy-Qwen audit compatibility, and 24-test suite are at `3bb12a8`.) (The managed remote does not accept tag refs; cite commit SHAs. Archiving a release snapshot, e.g. via Zenodo, is recommended before submission.)

Artifacts.

- Narrative and index of all experiments: `RESEARCH_NARRATIVE.md`.
- Preregistrations (frozen before data): `experiments/PREREG_*.md`, with item hashes, seeds, contrasts, thresholds, and missingness policy.
- Frozen items + oracles: `experiments/ccc_code_items.py` + `ccc_code_runner.py` (sandboxed unit-test grader); `experiments/ccc_sql_items.py` (SQLite fixtures, queries, canonicaliser, gold signature). Both self-verify and are hash-pinned; gold was checked byte-identical across CPython 3.10–3.13 and, for SQL, across SQLite versions via the run-time gold-signature gate.
- Run adapters: `experiments/run_ccc_codedomain.py`, `experiments/run_ccc_sql.py` (fixed concurrent worker pool, single writer; concurrency invariants verified offline and on live evidence).
- Unified per-domain findings: `experiments/results/FINDINGS_ccc_codedomain.md`, `.../FINDINGS_ccc_sql.md`.
- Frontier findings and independent audit: `experiments/results/FINDINGS_ccc_frontier.md` and `.../ccc_frontier_v3_audit.txt`.
- Frontier Phase 2 findings and audit: `experiments/results/FINDINGS_ccc_frontier_phase2.md` and `.../ccc_frontier_p2_audit.txt`.

- OpenRouter synthesis: `experiments/results/FINDINGS_ccc_openrouter_openweight.md`; model-specific raw evidence, findings, completion metadata, and audits are under `experiments/results/ccc_openrouter_*`.
- Offline audit logs (independent recomputation from raw rows): `.../ccc_code_offline_audit.txt` and `.../ccc_sql_offline_audit.txt`.
- Evidence: streamed `_obs_.jsonl`, `_prompts_.jsonl`, `_solver_.jsonl`, `_meta_.json` per stage; the `_meta_` files record seeds, hashes, oracle versions (e.g. SQLite 3.50.4 on the run machine), worker count, run date, and model aliases. Every reported estimate is auditable back to a stored transcript.

Endpoint provenance. Runs were executed 2026-07 through **OpenRouter**; model identifiers are OpenRouter slugs, i.e. provider aliases whose backing may change. Exact routing/version metadata beyond the alias is limited to what the router returned and is recorded in the `_meta_` files; readers reproducing the study should record their own endpoint dates and any routing metadata their provider exposes.

Integrity summary. Across the six cost-tier domain-stage runs: one row per intended cell, at most one attempt per cell, zero duplicate successful cells, zero order-index/cell mismatches under concurrency, all prompts resolvable in the manifests, factor-correlated missingness disclosed and fail-closed, and all headline contrasts recomputed by the independent pipeline. Frontier v3 contains all 5,376 intended rows and zero duplicate successes; its independent recomputation matches the reported estimates. Fourteen worker-level failures contain no usable provider choice trace, and Fable's injection-skewed filter blocks are explicitly reported as unmeasurable rather than safe. Frontier Phase 2 contains all 2,496 intended rows and zero duplicate successes; its independent recomputation applies the frozen balance rule, including the Gemini-SQL downgrade. Its shared metadata overwrite is disclosed as an audit exception. The four OpenRouter full arms each contain 1,344 unique successful cells, complete prompt manifests, fixed endpoint identities, and independently reproduced estimates. MiniMax and GLM additionally have 1,344/1,344 strict terminal JSON responses and zero observed reasoning tokens.

AI use in producing these experiments and analyses is described in "Author contributions and AI use".